

Minimum Deterministic Finite Automaton

SAT Modeling and Solution

David Garcia Tejeda

May 2025

Abstract

We present a SAT-based method for solving the Minimum Consistent Finite Automaton problem: finding the smallest deterministic finite automaton (DFA) that accepts all positive and rejects all negative binary sequences in a given labeled sample. The encoding, variables, constraints, and solving strategy are detailed.

1 Problem Description

A deterministic finite automaton (DFA) is defined by a set of states, an initial state, transitions for each input bit (0 or 1), and a set of accepting states. Given two disjoint sets of binary sequences A (accepted) and R (rejected), the task is to find a minimal-state DFA consistent with both.

A DFA accepts a sequence if, after processing all bits starting from the initial state, it reaches an accepting state; otherwise, it rejects it.

2 SAT Modeling

The problem is encoded into propositional satisfiability (SAT) as follows.

2.1 Variables

Let n be the number of states in the candidate DFA.

We define the following boolean variables:

- $T_{i,b,j}$: True if the DFA transitions from state i to state j on input bit $b \in \{0, 1\}$.
- A_i : True if state i is an accepting state.
- $S_{w,p,i}$: True if after reading the first p bits of word w , the DFA is in state i .

2.2 Constraints

1. **Deterministic Transitions:** For each state i and input bit b , exactly one $T_{i,b,j}$ is true for $j \in \{0, \dots, n-1\}$:

$$\sum_{j=0}^{n-1} T_{i,b,j} = 1$$

2. **Initial State:** For each word w , the empty prefix $p = 0$ leads to the initial state 0:

$$S_{w,0,0} = \text{true}, \quad S_{w,0,i} = \text{false} \quad \forall i \neq 0$$

3. **Unique State per Prefix:** For each word w and prefix length p , the DFA is in exactly one state:

$$\sum_{i=0}^{n-1} S_{w,p,i} = 1$$

4. **Transition Consistency:** For each word w , prefix length p , bit $b = w[p]$, and states i, j :

$$S_{w,p,i} \wedge T_{i,b,j} \implies S_{w,p+1,j}$$

5. **Acceptance and Rejection:** For each accepted word w , at least one accepting state is active after reading all bits:

$$\bigvee_{i=0}^{n-1} (S_{w,|w|,i} \wedge A_i)$$

For each rejected word w , at least one non-accepting state is active after reading all bits:

$$\bigvee_{i=0}^{n-1} (S_{w,|w|,i} \wedge \neg A_i)$$

2.3 Encoding Details

All the above logical constraints are translated to Conjunctive Normal Form (CNF) using standard techniques, such as:

- Exactly-one constraints are encoded with pairwise mutual exclusion or more efficient encodings such as sequential counters.
- Implications are converted to CNF clauses.
- State variables $S_{w,i}$ are introduced for all prefixes of sequences in $A \cup R$.

3 Implementation and Solving

The solver is implemented in C++ as a single-binary project with Kissat as its only dependency. It reads the problem instance from `stdin`, dynamically encodes the constraints into CNF based on the input sample and current DFA size n , and calls the SAT solver.

Constraint encoding is performed incrementally: the system generates only the clauses required for the current n , leveraging the SAT solver’s internal mechanisms (e.g., clause learning, propagation) to manage the problem’s combinatorial complexity.

The search proceeds by incrementing n from 1 upward until a satisfying assignment is found, which guarantees minimality. The resulting model is decoded to reconstruct a valid DFA, printed in the specified format.

4 Remarks and Observations

- **Dynamic Constraint Encoding:** The model supports generating alternative versions of the `atMostOne` constraint encoding dynamically, adapting heuristics such as the number of involved variables. Constraints are produced on demand for each candidate size n , avoiding the overhead of monolithic encodings and enabling scalable, efficient exploration of the search space.

- **Solver-Aided Complexity Management:** By offloading the combinatorial reasoning to the SAT solver, the system benefits from efficient clause learning, propagation, and conflict analysis.
- **Minimal DFA Guarantee:** The incremental search over increasing n ensures the first solution found is minimal by construction.
- **Resource Efficiency:** Only necessary state and transition variables are introduced based on actual sequence prefixes, minimizing redundant logic.
- **Timeout Handling:** A solver-level timeout (e.g., 60 seconds) ensures practical runtime bounds without modifying core logic.
- **Extensibility:** The modular encoding strategy allows future integration of richer constraints (e.g., transition symmetry, forbidden patterns) without altering the solving backbone.
- **Future Work - Search Optimization:** A more efficient search can be implemented by first performing an exponential (doubling) search to find an upper bound where a solution exists, and then switching to a binary (dichotomous) search between the last unsuccessful and first successful bounds. This hybrid approach reduces the number of solver invocations to roughly $O(\log n)$, improving overall performance compared to a linear search. It can be tested that trying to find a suboptimal solution (a DFA much larger than the minimum) is not a very costly operation as expected, since although we are generating more variables the search space is more rich in possible solutions and their isomorphisms.

5 Conclusion

The SAT-based modeling provides a clear and flexible framework for solving the Minimum Consistent Finite Automaton problem. By leveraging state-of-the-art SAT solvers, this approach can find minimal DFAs consistent with given samples, supporting the combinatorial problem solving objectives.

6 Final remarks

All code can be found at https://github.com/DavidNexuss/cps_3 or in the attached zip.